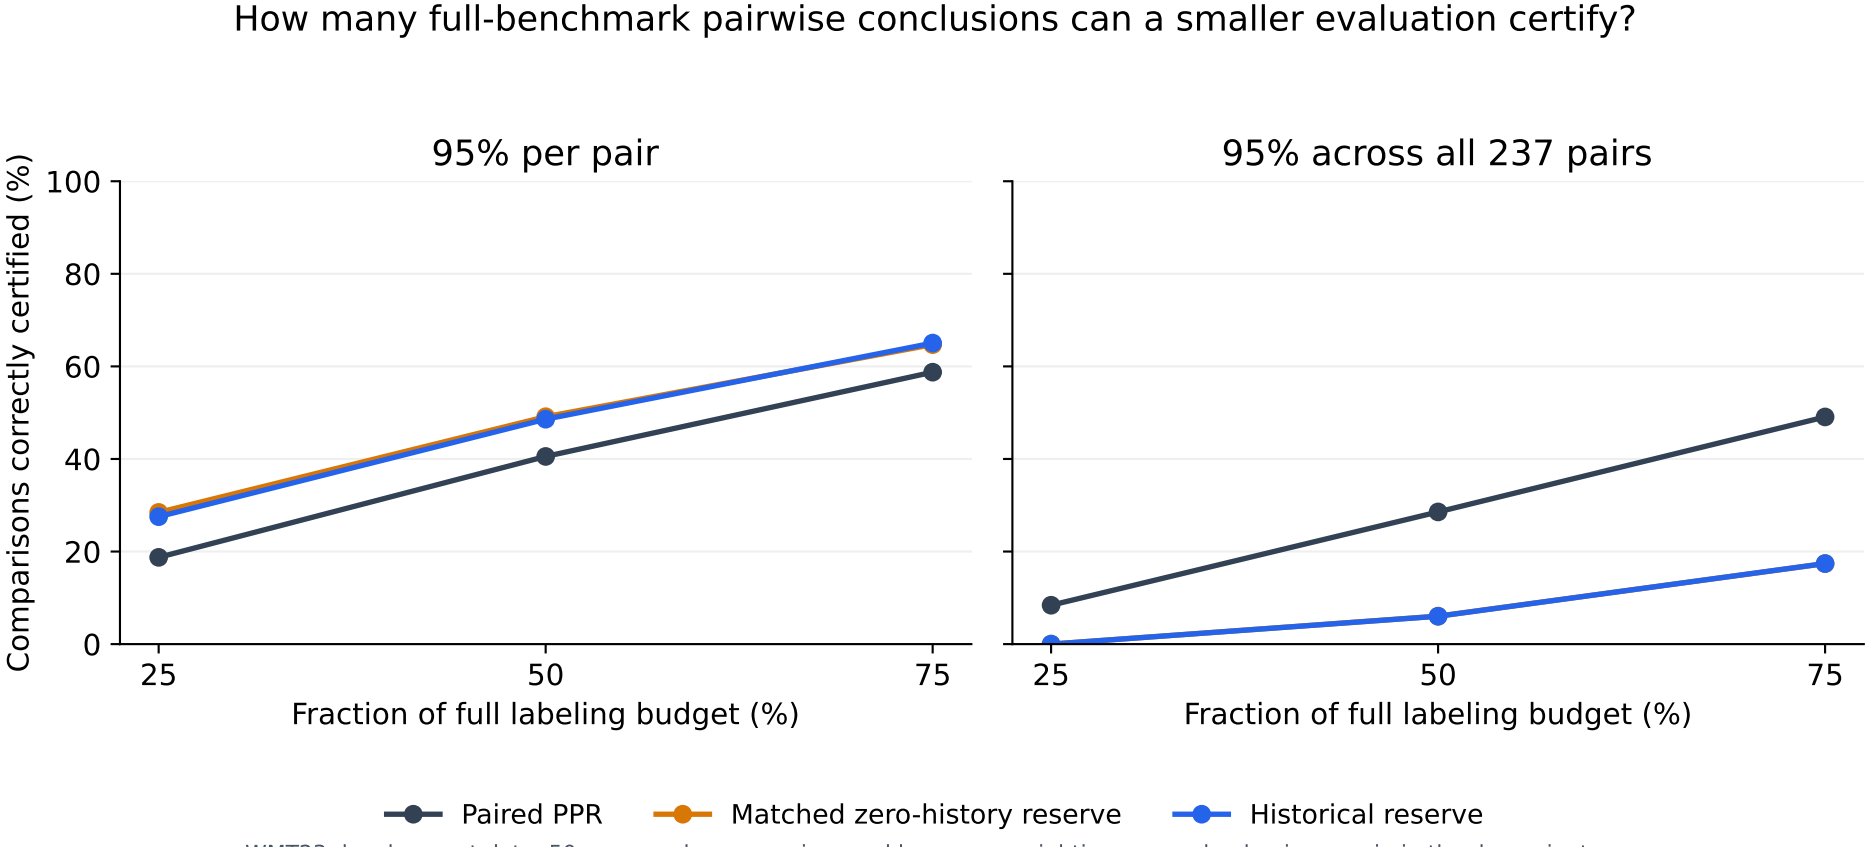
How many full-benchmark pairwise conclusions can a smaller evaluation certify?
95% per pair
95% across all 237 pairs
Comparisons correctly certified (%)
100
80
60
40
20
0
25
50
75
Fraction of full labeling budget (%)
Paired PPR
Matched zero-history reserve
Historical reserve
WMT23 development data; 50 query orders per pair; equal language weighting; unresolved pairs remain in the denominator.
Family-wide reserve curves have a component-cap limitation (see CAP_LIMITATION.md). Error settings apply separately to each method.